

OPEN SOURCE CLOUD COMPUTING FORUM

February 10, 2010

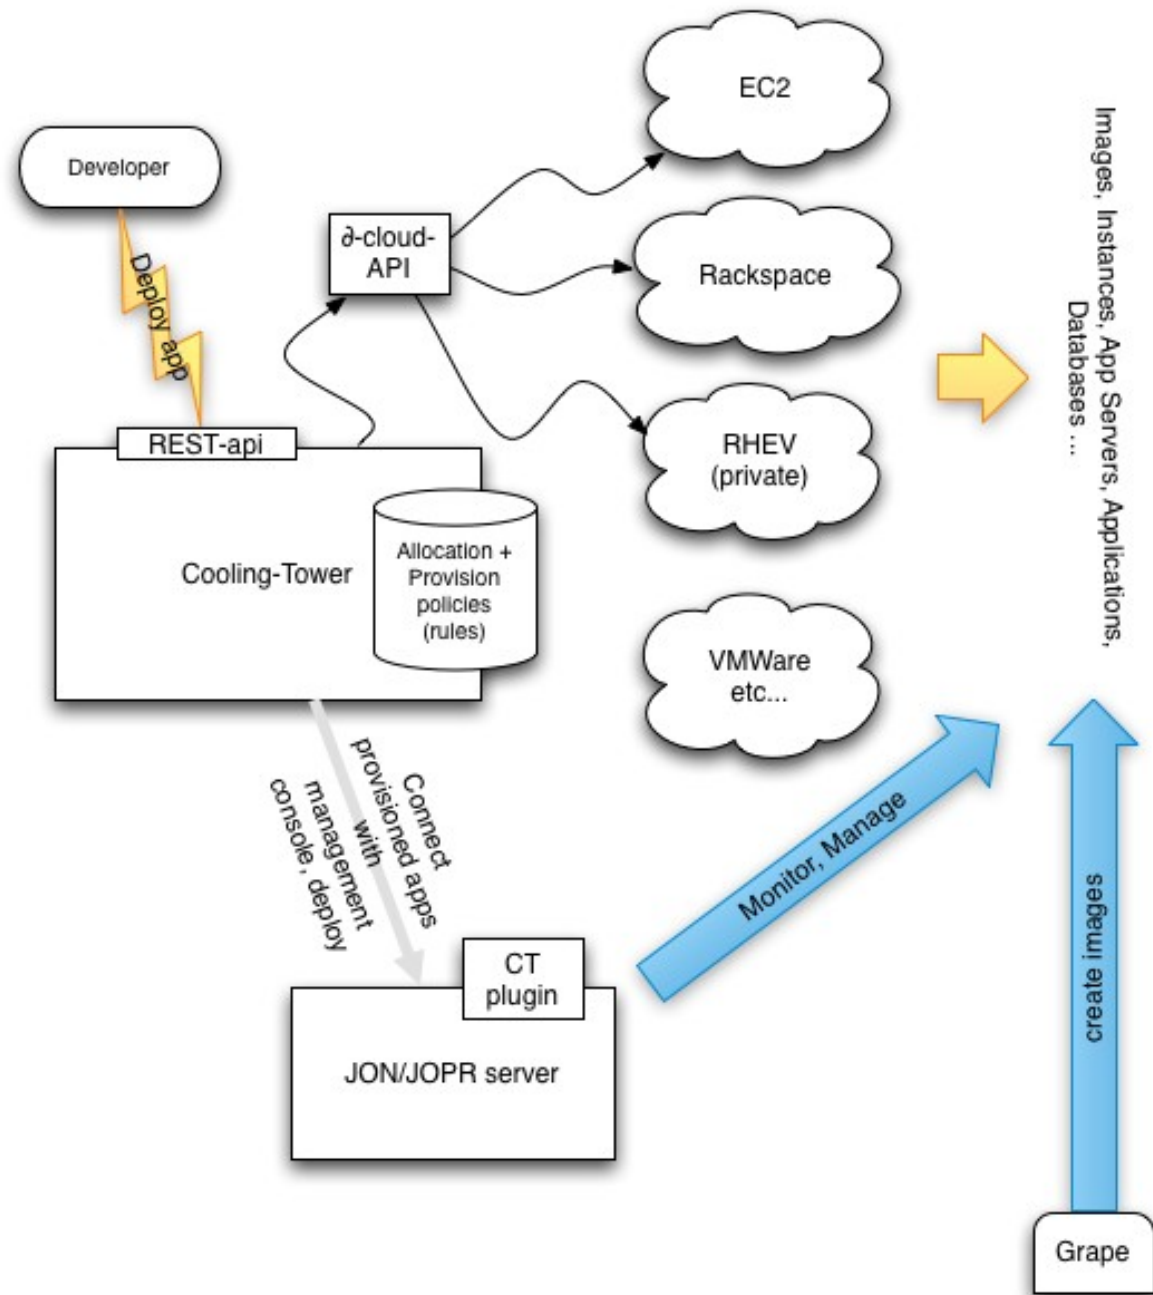
CoolingTower

Michael Neale
Senior Engineer, Red Hat
(JBoss)

CoolingTower: Turnkey PaaS

- What is CoolingTower
 - A platform-as-a-service for Java Server side apps (Servlet/EE container) – completely self service
 - Application centric (it is a PaaS) not server/infrastructure centric
 - Is a project of utilities and aggressive automation of existing (or soon to be existing) components
 - Runs as a server component with RESTful API
 - Policy driven
 - Driven by RESTful apis, multiple front ends
 - Using open API's such as deltacloud for multiple IaaS platforms

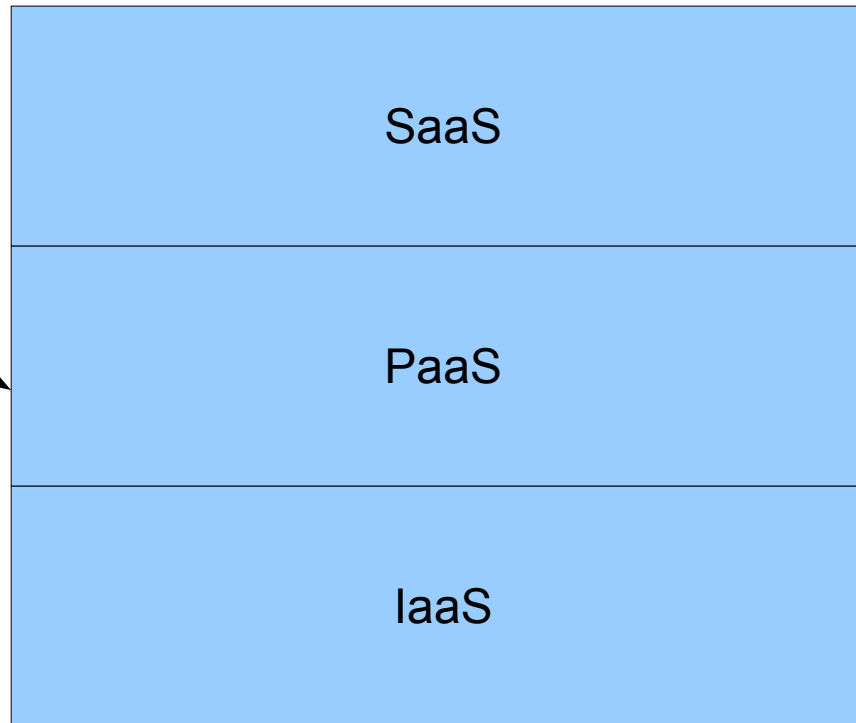
CoolingTower:



PaaS refresher...

“off the shelf”, eg
Expenses, Timesheets,
CRM

Apps:
hr.yourcloud.com,
app2.yourcloud.com



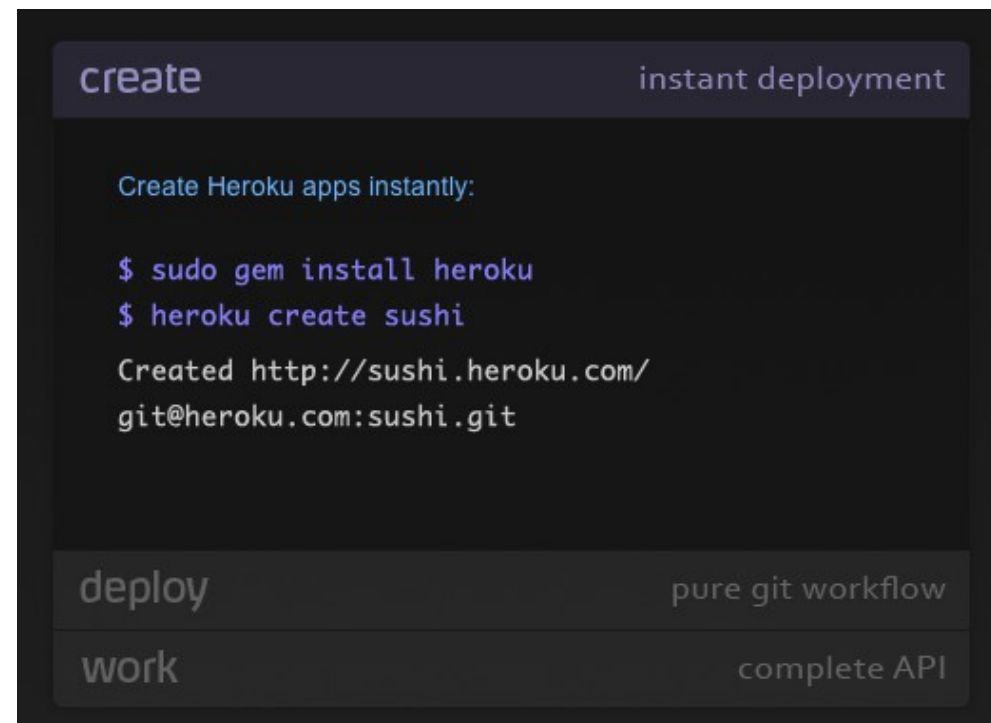
server1west, db02 etc.

Motivation

- Many clouds are infrastructure and operations centric
- A space exists for application and developer centric clouds
 - Yet adhering to organisational policies, workflows for application lifecycles etc.
 - Automation !
- Build on standards !
 - emerging for clouds, and for apps

Public PaaS already in use

- Heroku.com
 - Ruby based



Google App engine

- Java or Python editions
 - Free or Pay (based on quotas)
 - Can run other languages via JVM



Run your web apps on Google's infrastructure.
Easy to build, easy to maintain, easy to scale.

What do they have in common?

- Stunning developer experience (ease, low/no barrier)
- Instant results (via a nice URL)
- Some ability to scale elastically *
- NO management of server resources *
- Infrastructure abstracted (automated) away
 - You don't manage/see servers
- Very good uptake for small/medium apps
- Very ATTRACTIVE to **developers**
 - **(who don't want to administer !)**

Limitations (not all bad)

- Often a subset of a “full stack”
 - Must stick within API
 - Locked to specific platform versions
 - Build apps to scale
- Run on public infrastructure (shared)
- Follow the conventions, live with with specific versions
- Don't have usual “enterprise” needs for promotion, QA environments etc

CoolingTower - motivations

- Use standards as much as possible
 - JavaEE, Servlet API (JVM allows many languages now)
 - Deltacloud for IaaS layer
- A PaaS cloud is a PaaS cloud: avoid lock in to IaaS layer
- Allow PaaS cloud to be built on private infrastructure, or public
- Address “enterprise concerns”
 - Use configurable deployment and management policies
 - Allow promotion through cloud “environments” (dev/test → QA → production etc).

Further motivations

- In enterprise, apps need to scale in many directions:
 - Up: lots of users
 - Deep: complex app with lots of calculations (but maybe not many users)
 - Many: many small departmental apps – come and go.
- Reality: many apps don't need to scale like GAE style apps do (but high availability is very nice if possible !)
- With JEE/JBossAS we get relatively easy app tier clustering

CoolingTower - components

- ***Deployment API***
 - RESTful api, Policy driven deployment
- ***Elastic scaler***
 - Policy driven horizontal scaling of clustered apps
- ***Server minimiser***
 - “densely” pack applications for development clusters
- ***Naming Service***
 - RESTful api for self service of application names
 - Authoritative only DNS
- ***Clients: Web, Eclipse, RESTful***

CoolingTower - deployment

- Applications packaged up in standard war/ear (optional extra meta data)
- Deploy via RESTful API

Usage of REST api (draft)

```
GET /cooling-tower -- get a list of links to something like the following:
POST /cooling-tower/applications --- post the war/app contents, will return a link to status
GET /cooling-tower/applications/your-app -- return the status of it (where it is running, if it is ready etc)
POST /cooling-tower/applications/your-app -- update the version
DELETE /cooling-tower/applications/your-app -- decommission the app (may result in some server savings)
GET /cooling-tower/applications -- list the apps available and status
```


CoolingTower - deployment

- On request:
 - service applies policy rules, finds a running instance, or starts instances as needed (using CirrAS to start a cluster up of app servers)
 - deploys application, runs promotion “tasks”
- Track status of applications via RESTful call
- Update an application easily, decommission
- Policy could enforce what certain developers can deploy to – notify people of new version
 - Workflow: new version, approval, promotion to QA, PROD etc.

Deployment - policies

- Currently uses deltacloud “model”
- Policies are declarative rules !
- Selection of appropriate image, or running instance

Deployment - policies

```
rule "Can't colocate - so need a new server"
when
    application: Application(canCoExist == false)
    not Instance(numberOfApps == 0,
                  spareMemory >= application.memory,
                  spareStorage >= application.disk)
then
    insert(new NewInstanceNeeded(application))
end

#
#TIP: you can lock it down to specific image types here if you like...
#
rule "Can colocate - see if we can fit it in somewhere"
when
    application: Application(canCoExist == true)
    instance: Instance(numberOfApps < 4,
                       spareMemory >= (application.memory * 1.2),
                       spareStorage >= (application.disk * 1.5) )
    #we want to make sure there is no other instance with less apps !
    not Instance(numberOfApps < instance.numberOfApps,
                  spareMemory >= (application.memory * 1.2),
                  spareStorage >= (application.disk * 1.5) )
    forall(Application(canCoExist == true) from instance.applications)
then
    insert(new Assignment(application, instance))
end
```

Declarative rules

- Like scripts, but far more powerful
- Combine data of cloud (images, instances) with application needs (meta data, other running applications) to reason over in rules
- Efficient use of cloud IaaS layer
- Rules, so customisable for a given cloud installation
- Cover:
 - Allocation for deployment, image selection
 - instance life cycle
 - application lifecycle (promotion)

Roadmap: deployment front ends

- Eclipse
 - Developers deploy from eclipse to development cloud (whence workflow for promotion can be initialised)
- Command line
 - Scripts, ant (all good RESTful interfaces should be able to be driven from the command line)
- Web Client
 - Web based PaaS console for creating/updating apps

Roadmap: data services

- Most practical applications require persistent data
 - NoSQL (Infinispan REST api – as well as cache)
 - But not for everyone
 - RDBMS – will need to offer a fixed set (Postgres likely, possibly MySQL in limited configuration)
 - Most JVM apps use ORM which provides RDBMS flexibility
 - TBD...

CoolingTower – Elastic Scaler

- Once again, use declarative policy rules
- Commonly done with scripts: monitor load, grow and shrink as needed (application tier mainly)
- Rules will apply that, as well as conform to policy
 - limiting resource load
 - prioritising for important apps
 - attempt to anticipate load and scale up ahead of time (user driven, or based on historical data patterns).

A note on app multi-tenancy

- Instances can run multiple application servers
- Application servers (1 process) can run multiple applications
- Need to have policies to decide what level of isolation is needed for each app.
- For even “lite” applications, can have 2 instances supporting it for HA reasons
 - So multi-tenanting often makes sense
- For development purposes, may want to multi tenant for efficiency.


CoolingTower – Server Minimiser

- Keep costs low (both in \$ and resources)
- Many IaaS “bill” per-instance
- Keep instance numbers as low as possible for applications
- Mainly important for development clusters
- Each process/app has a footprint of memory, CPU, disk needs, an instance has finite amount of each
- Resembles the “bin packing problem”

Minimiser

- We have a “constraint solver” (called Drools Planner)
- Declare rules to describe the hard and soft constraints
- Try multiple configurations of applications per instance to result in a solution with lower number of instances needed
- Could be run periodically (in the case of many apps coming and going)

Minimiser



Minimiser: roadmap

- Very experimental
 - Most likely of use for public IaaS where cost is clear. For private clouds, not as urgent
- May be important longer term
 - Minimise hardware foot print
 - Will need to take into account data we don't yet easily have (eg energy usage) or understand

CoolingTower: Naming Service

- The “forgotten” cloud service
 - <http://www.rackspacecloud.com/blog/2009/06/04/dns-the-overlooked-cloud-service/>
- In a PaaS cloud, typically have many apps:
 - `hr.yourcloud.com`, `incentive.yourcloud.com` are nice
 - Web “entry-point” may move around
 - Apps can come and go
 - Eg EC2 “ugly” names, or rackspace IPs
- You may want to refer to “cloud services” via URLs, eg caching, vertical services etc

Naming Service

- Typical process
 - When setting up a cloud, buy domain name
 - Register CoolingTower/PaaS clouds DNS servers with registrar you use (own DNS gives control)
 - CoolingTower does the rest
 - DNS can be error prone, so automate !

Naming Service – a RESTful api for DNS

- Register a domain
 - POST /cooling-tower/naming name=yourcloud.com
- Set/update address of subdomain (or default)
 - POST ../yourdomain.com
subdomain=applicationName&address=1.2.3.4
 - PUT ../yourdomain.com/www address=newAddress...
- (IP or domain name, either works)
- Many more REST verbs apply for modifying, removing etc.
- CoolingTower will call this when creating an instance etc.

Naming – RESTful DNS

- CoolingTower stores DNS records in standard zone files
- Can rsync to BIND servers (so the “master” is hidden), or use built in authoritative server (which is updated instantly with changes)
- Use very short TTLs to allow services to move around in the cloud
- For both internal and external services

Work to be done..

- Defining “cloud descriptor” for meta data needed to deploy
 - Include automating of data/migration scripts when “promoting” an application
- Defining common RDBMS services that cover enough use-cases
- Integration with CirrAS for deployment of applications into clustered App Servers.
- Capturing more policies/samples in rules
- Much more...

Community: Getting involved

- Subproject of StormGrind @ JBoss
 - jboss.org/stormgrind/projects/coolingtower
 - Code currently on [github/michaelneale/cooling-tower](https://github.com/michaelneale/cooling-tower)
 - wiki.github.com/michaelneale/cooling-tower
 - (will migrate to jboss.org)
 - IRC: [#stormgrind](https://irc.freenode.net) on irc.freenode.net